

In [5]:

```
#importing the required Packages
import pandas as pd
from textblob import TextBlob
import re
import matplotlib.pyplot as plt
import seaborn
import itertools
import string
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from wordcloud import WordCloud
from wordcloud import STOPWORDS
from sklearn import ensemble
from sklearn import tree
from sklearn import metrics
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC, LinearSVC
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.pipeline import Pipeline
from sklearn.cross_validation import cross_val_score, train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.preprocessing import Binarizer, StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
from nltk.sentiment.util import mark_negation
from matplotlib import pylab
pylab.rcParams['figure.figsize'] = (15, 9)

import warnings
warnings.filterwarnings("ignore")
```

In [6]:

```
##### Defining the required Defined Functions

#%% Functions
def clean_tweet(tweet):
    '''
    Utility function to clean tweet text by removing links, special characters
    using simple regex statements.
    '''
    return ' '.join(re.sub("(@[A-Za-z0-9]+)|([^0-9A-Za-z \t])|(\w+:\/\/\S+)", " ", str(tweet)).
split())

def get_tweet_sentiment(tweet):
    '''
    Utility function to classify sentiment of passed tweet
    using textblob's sentiment method
    '''
    # create TextBlob object of passed tweet text
    analysis = TextBlob(clean_tweet(tweet))
    # set sentiment
    if analysis.sentiment.polarity > 0:
        return 'positive'
    elif analysis.sentiment.polarity == 0:
        return 'neutral'
    else:
        return 'negative'

def tweet_clean(df):
    temp_df = df.copy()
    # Remove hyperlinks
    temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('https?:\/\/\.\.*\/\w*', '', regex=True)
    # Remove citations
    temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('\@\w*', '', regex=True)
    # Remove tickers
```

```

temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('\$\w*', '', regex=True)
# Remove punctuation
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('[' + string.punctuation + ']+', '', regex=True)
# Remove quotes
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('&*[amp]*\;|gt+', '', regex=True)
# Remove RT
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('RT', '', regex=True)
# Remove linebreak, tab, return
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('[\n\t\r]+', ' ', regex=True)
# Remove via with blank
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('via\s', '', regex=True)
# Remove multiple whitespace
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('\s+\s+', ' ', regex=True)
# Remove multiple whitespace
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('\s+\s+', ' ', regex=True)
# Remove HashTags
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('\#+[\w_]+[\w\'_\-]*[\w_]+', ' ', regex=True)
# Remove Smileys
temp_df.loc[:, "Text"] = temp_df.loc[:, "Text"].replace('[:]=+(|o|O| )+[D\)]\))+[\(\[]+[pP]+[doO/\]]+[\(\[]+(\^_\^|)', ' ', regex=True)
# Remove empty rows
temp_df = temp_df.dropna()
return temp_df

def regularExpression(textToFilter):
    filteredTweet = []
    retweetPattern = 'RT|@RT'
    urlPattern = 'https://[a-zA-Z0-9+&@#/%?~_!|:,.;]*'

    for textLine in textToFilter:
        tweet = re.sub(retweetPattern, '', textLine)
        tweet = re.sub(urlPattern, '', tweet)
        filteredTweet.append(tweet)
    return filteredTweet

def nltkTokenizer(textToTokenize):
    filteredSentence = []
    usersPattern = re.compile('@[a-zA-Z0-9]*', re.UNICODE)
    hashtagPattern = re.compile('#[a-zA-Z0-9]*', re.UNICODE)
    stop_words = stopwords.words('english')

    for textLine in textToTokenize:
        words = re.sub(usersPattern, '', textLine)
        words = re.sub(hashtagPattern, '', words)
        words = word_tokenize(words)
        for w in words:
            if w not in stop_words and w not in '@' and w not in '#':
                filteredSentence.append(w)
    return filteredSentence

def tweet_to_words(raw_tweet):
    tweet = ''.join(c for c in raw_tweet if c not in string.punctuation)
    tweet = re.sub('((www\S+)|(http\S+))', 'urlsite', tweet)
    tweet = re.sub(r'\d+', 'contnum', tweet)
    tweet = re.sub(' +', ' ', tweet)
    words = tweet.lower().split()
    stops = set(stopwords.words("english"))

    meaningful_words = [w for w in words if w not in stops]
    return " ".join(meaningful_words)

def users(tweet):
    user = []
    usersPattern = re.compile('@[a-zA-Z0-9]*', re.UNICODE)

    for t in tweet:
        u = re.findall(usersPattern, t)
        user.append(u)
    return user

def split_into_tokens(Text):
    return TextBlob(Text).words

def split_into_lemmas(Text):
    Text = Text.lower()

```

```
words = TextBlob(Text).words
# for each word, take its "base form" = lemma
return [word.lemma for word in words]
```

In [7]:

```
file_path = "C:\\Users\\91953\\Desktop\\Analytic Labs\\Thesis and Course\\Dublin Business School\\Subham\\Thesis\\Data\\All_Tweets.csv"
```

In [8]:

```
tweets = pd.read_csv(file_path)

print(tweets.shape[0])

tweet = tweets["Text"]

cleaed_Tweet = clean_tweet(tweet)

Polarity = []
for tweet in tweets['Text']:
    Polarity.append(get_tweet_sentiment(tweet))

tweet = tweets['Text']
data_sent = {'Text': tweet, 'Polarity': Polarity}
tweet_data = pd.DataFrame(data=data_sent)

print("=====")
print("Printing the Head of the Tweets: =====")
print(tweet_data.head())
print("=====")
```

```
12000
=====
Printing the Head of the Tweets: =====
                                     Text  Polarity
0  89% of the cases reported today are asymptomat...  positive
1  89% of the cases reported today are asymptomat...  positive
2  @BarbaraGirouard @JamesE2020 @FriendsOScience ...  neutral
3  89% of today's record high Covid-19 cases in K...  positive
4  89% of the cases reported today are asymptomat...  positive
=====
```

In [10]:


```
print("Plotting the TextBlob Sentiments: =====")
Index = [1,2,3]
print(tweet_data.Polarity.value_counts())
plt.bar(Index,tweet_data.Polarity.value_counts())
plt.xticks(Index,['neutral','positive','negative'],rotation=45)
plt.ylabel('Number of Posts')
plt.xlabel('Sentiment expressed in Posts')
```

```
Plotting the TextBlob Sentiments: =====
neutral      5247
positive     4198
negative     2555
Name: Polarity, dtype: int64
```

Out[10]:

Text(0.5,0,'Sentiment expressed in Posts')





In [11]:

```
polar = pd.DataFrame()
n = int(len(tweet))
sen = []
for i in range(n):
    blob = TextBlob(str(tweet[i]))
    k = blob.sentiment.polarity
    sen.append(k)

polar['polarity'] = sen
print("Printing the Polar Data Head: =====")
print(polar.head())
```

```
Printing the Polar Data Head: =====
   polarity
0  0.219444
1  0.219444
2  0.000000
3  0.086010
4  0.219444
```


In [12]:



```
polar.hist()
polar.plot.line(y='polarity',figsize=(18,9))

tweet_data.groupby('Polarity').describe()

tweet_data['length'] = tweet_data['Text'].map(lambda text: len(text))
print("Printing the head of Tweets Data Newly Formed: =====")
print(tweet_data.head())
print("=====")
```

```
Printing the head of Tweets Data Newly Formed: =====
   Text  Polarity  length
0  89% of the cases reported today are asymptomat...  positive    162
1  89% of the cases reported today are asymptomat...  positive    162
2  @BarbaraGirouard @JamesE2020 @FriendsOScience ...   neutral    589
3  89% of today's record high Covid-19 cases in K...  positive    297
4  89% of the cases reported today are asymptomat...  positive    162
=====
```






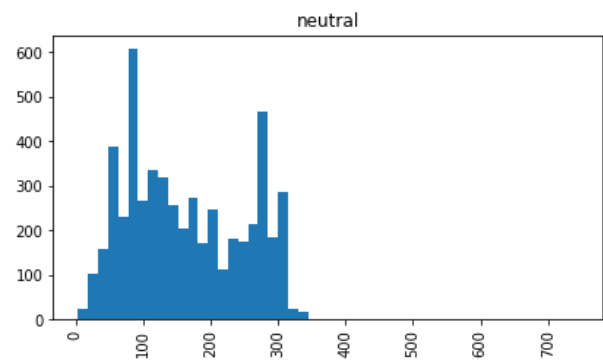


In [13]:

```
print("Histogram Plot of the Frequency: =====")
tweet_data.length.plot(bins=20, kind='hist')

tweet_data.hist(column='length', by='Polarity', bins=50)
print("Printing the Shape of the Tweets: =====")
print(tweet_data.shape)
```

```
Histogram Plot of the Frequency: =====
Printing the Shape of the Tweets: =====
(12000, 3)
```





In [14]:

```
tweets_texts = tweet_data["Text"].tolist()
stop_words=stopwords.words('english')
english_vocab = set(w.lower() for w in nltk.corpus.words.words())

# Get all the hashtag words that has "#"
hashtags = ""
for line in tweets:
    words = line.split()
    for w in words:
        if w.startswith("#"):
            hashtags += w + " "

# Get all the hashtags in a list
hashtags_list = re.findall(r"#(\w+)", hashtags)

print("Plotting the Hashtag WordCloud: =====")
try:
    # Set the figure-size
    plt.figure(figsize= (20,10))
    wordcloud = WordCloud(
        stopwords=STOPWORDS,
        background_color='white',
        width=3000,
        height=2000
    ).generate(str(hashtags_list))

    plt.figure(1,figsize=(20, 20))
    plt.imshow(wordcloud)
```

```

plt.axis('off')
plt.savefig('./twitter_wordcloud.png', dpi=300)
plt.show()
except:
    print("No HashTags Found!! Please Verify the length")

print(len(hashtags_list))
print("=====")

```

```

Plotting the Hashtag WordCloud: =====
No HashTags Found!! Please Verify the length
0
=====

```

<Figure size 1440x720 with 0 Axes>

In [15]:

```

###
# Text contains 'RT' for every retweet and url references
# We want to remover 'RT' and URL

filteredTweet = regularExpression(tweet_data.Text)
filteredSentence = nltkTokenizer(filteredTweet)

hashtagList = list(itertools.chain.from_iterable(hashtags_list))
hashtagCount = {}

for h in hashtagList:
    if h in hashtagCount:
        hashtagCount[h] +=1
    else:
        hashtagCount[h] = 1

# Extracting hastags that occurs more than 1000 times

hashtagCount = { k : v for k,v in hashtagCount.items() if v >10}
name = [k for k in hashtagCount if k ]
value = [v for v in hashtagCount.values()]

```

In [16]:

```

print("Extracting the User words used in the Tweets: =====")
user = users(filteredTweet)

str1 = ' '.join(str(e) for e in user)

plt.figure(figsize= (20,20))


wordcloud = WordCloud(width= 3000,height=
2000,background_color='white',max_words=30).generate(str1)
plt.imshow(wordcloud)
plt.title('Most mentioned users')
plt.axis("off")
plt.show()
print("=====")

```

```

Extracting the User words used in the Tweets: =====

```




In [18]:

```
pd.Series(tweet_data["Polarity"]).value_counts().plot(kind="pie", colors=colors,
    labels=["negative", "neutral", "positive"], explode=[0.05, 0.02, 0.04],
    shadow=True, autopct='%2f', fontsize=12, figsize=(12, 12), title = "Total Tweets for Each Sentimen
t")

df=tweet_data[tweet_data['Polarity']=='negative']
words = ' '.join(df['Text'])
cleaned_word = " ".join([word for word in words.split()
    if 'http' not in word
    and not word.startswith('@')
    and word != 'RT'])

wordcloud = WordCloud(stopwords=STOPWORDS,
    background_color='white',
    width=3000,
    height=2500
).generate(cleaned_word)

plt.figure(1, figsize=(20, 20))
plt.imshow(wordcloud)
plt.axis('off')
plt.savefig('./negative_tweet_wordcloud.png', dpi=300)
plt.show()
```



In [19]:

```
df=tweet_data[tweet_data['Polarity']=='positive']
words = ' '.join(df['Text'])
cleaned_word = " ".join([word for word in words.split()
    if 'http' not in word
    and not word.startswith('@')
    and word != 'RT'])
```


[illegible]

```
df=tweet_data[tweet_data['Polarity']=='neutral']
words = ' '.join(df['Text'])
cleaned_word = " ".join([word for word in words.split()
                           if 'http' not in word
                           and not word.startswith('@')
                           and word != 'RT']])

wordcloud = WordCloud(stopwords=STOPWORDS,
                       background_color='white',
                       width=3000,
                       height=2500
                       ).generate(cleaned_word)

plt.figure(1,figsize=(20, 20))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
print("=====")
```

honneur du Y viru infection wishal tip koronaviruss fortes chaleurs rds et



In [22]:

```
print("=====")
print("Starting the Machine Learning Analysis: =====")
print("=====")

tweet_data['clean_tweet']=tweet_data['Text'].apply(lambda x: tweet_to_words(x))
train,test = train_test_split(tweet_data,test_size=0.33,random_state=0)

train_clean_tweet=[]
for tweet in train['clean_tweet']:
    train_clean_tweet.append(tweet)
test_clean_tweet=[]
for tweet in test['clean_tweet']:
    test_clean_tweet.append(tweet)
print("Using the COUNTVECTORIZER: =====")
v = CountVectorizer(analyzer = "word")
train_features= v.fit_transform(train_clean_tweet)
test_features=v.transform(test_clean_tweet)

Classifiers = [
    LogisticRegression(C=0.001,multi_class='multinomial',max_iter=10,solver='sag', tol=1e-1),

    RandomForestClassifier(n_estimators=200, bootstrap=True, class_weight=None, criterion='gini',
        max_depth=50, max_features='auto', max_leaf_nodes=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, n_jobs= -1,
        oob_score=False, random_state=10),

    AdaBoostClassifier(n_estimators=100, random_state=10),

    BernoulliNB(),

    MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True),

    KNeighborsClassifier(algorithm='auto', metric='minkowski',
        metric_params=None, n_neighbors=2, p=2,
        weights='uniform'),

    tree.DecisionTreeClassifier(),

    ensemble.ExtraTreesClassifier(n_estimators=100,
        max_features= 50,
        criterion= 'entropy'),

    ensemble.GradientBoostingClassifier(criterion='friedman_mse', init=None,
        learning_rate=0.001,n_estimators=50,presort='auto', random_state=None, verbose = 0)]
```

```

dense_features=train_features.toarray()
dense_test= test_features.toarray()
Accuracy=[]
Model=[]
print("Entering into the Classifiers: =====")
for classifier in Classifiers:
    try:
        fit = classifier.fit(train_features,train['Polarity'])
        pred = fit.predict(test_features)
    except Exception:
        fit = classifier.fit(dense_features,train['Polarity'])
        pred = fit.predict(dense_test)
    accuracy = accuracy_score(pred,test['Polarity'])
    print("=====")
    print("*****")
    print('Accuracy of '+classifier.__class__.__name__+'is '+str(accuracy))
    print("*****")
    print("=====")
    Accuracy.append(accuracy)
    print("Classification report for classifier %s:\n%s\n"
          % (classifier, metrics.classification_report(test['Polarity'], pred)))
    cm = metrics.confusion_matrix(test['Polarity'], pred)
    print("-----")
    print("Confusion matrix:\n%s" % cm)
    print("-----")
    Model.append(classifier.__class__.__name__)
    print("*****")
    print("=====")

```

=====

Starting the Machine Learning Analysis: =====

=====

Using the COUNTVECTORIZER: =====

Entering into the Classifiers: =====

=====

Accuracy of LogisticRegressionis 0.7080808080808081

=====

Classification report for classifier LogisticRegression(C=0.001, class_weight=None, dual=False, fit_intercept=True,

intercept_scaling=1, max_iter=10, multi_class='multinomial',
n_jobs=1, penalty='l2', random_state=None, solver='sag', tol=0.1,
verbose=0, warm_start=False):

	precision	recall	f1-score	support
negative	0.90	0.29	0.44	857
neutral	0.86	0.76	0.81	1757
positive	0.57	0.90	0.70	1346
avg / total	0.77	0.71	0.69	3960

Confusion matrix:

```

[[ 251  98  508]
 [   6 1341  410]
 [   21  113 1212]]

```

=====

Accuracy of RandomForestClassifieris 0.8113636363636364

=====

Classification report for classifier RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',

max_depth=50, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=-1,
oob_score=False, random_state=10, verbose=0, warm_start=False):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

negative	1.00	0.47	0.64	857
neutral	0.81	0.95	0.87	1757
positive	0.76	0.85	0.80	1346
avg / total	0.83	0.81	0.80	3960

Confusion matrix:

```
[[ 405  189  263]
 [    0 1665   92]
 [    1  202 1143]]
```


=====

Accuracy of AdaBoostClassifier is 0.8037878787878788

Classification report for classifier AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None, learning_rate=1.0, n_estimators=100, random_state=10):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

negative	0.81	0.63	0.71	857
neutral	0.78	0.96	0.86	1757
positive	0.85	0.71	0.77	1346
avg / total	0.81	0.80	0.80	3960

Confusion matrix:

```
[[ 544  184  129]
 [   25 1690   42]
 [  100  297  949]]
```


=====

Accuracy of BernoulliNB is 0.7863636363636364

Classification report for classifier BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

negative	0.96	0.50	0.65	857
neutral	0.93	0.81	0.87	1757
positive	0.64	0.94	0.76	1346
avg / total	0.84	0.79	0.78	3960

Confusion matrix:

```
[[ 425   40  392]
 [    5 1428  324]
 [   13   72 1261]]
```


=====

Accuracy of MultinomialNB is 0.7638888888888888

Classification report for classifier MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

negative	0.84	0.59	0.70	857
neutral	0.94	0.72	0.82	1757
positive	0.62	0.93	0.74	1346
avg / total	0.81	0.76	0.77	3960

```
-----
Confusion matrix:
[[ 508   21  328]
 [   53 1268  436]
 [   43   54 1249]]
-----
```

```
*****
=====
```

```
*****
Accuracy of KNeighborsClassifieris 0.6941919191919191
*****
```

```
=====
```

```
Classification report for classifier KNeighborsClassifier(algorithm='auto', leaf_size=30,
metric='minkowski',
```

```
metric_params=None, n_jobs=1, n_neighbors=2, p=2,
weights='uniform'):
```

	precision	recall	f1-score	support
negative	0.94	0.55	0.70	857
neutral	0.60	0.99	0.75	1757
positive	0.97	0.39	0.56	1346
avg / total	0.80	0.69	0.67	3960

```
-----
Confusion matrix:
[[ 474  370   13]
 [    9 1745    3]
 [   22  794  530]]
-----
```

```
*****
=====
```

```
*****
Accuracy of DecisionTreeClassifieris 0.825
*****
```

```
=====
```

```
Classification report for classifier DecisionTreeClassifier(class_weight=None, criterion='gini', m
ax_depth=None,
```

```
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best'):
```

	precision	recall	f1-score	support
negative	0.81	0.67	0.73	857
neutral	0.86	0.91	0.88	1757
positive	0.79	0.81	0.80	1346
avg / total	0.82	0.82	0.82	3960

```
-----
Confusion matrix:
[[ 574  106  177]
 [   39 1600  118]
 [   97  156 1093]]
-----
```

```
*****
=====
```

```
*****
Accuracy of ExtraTreesClassifieris 0.8550505050505051
*****
```

```
=====
```

```
Classification report for classifier ExtraTreesClassifier(bootstrap=False, class_weight=None,
criterion='entropy',
```

```
max_depth=None, max_features=50, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
oob_score=False, random_state=None, verbose=0, warm_start=False):
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

negative	0.96	0.61	0.74	857
neutral	0.87	0.96	0.91	1757
positive	0.80	0.88	0.84	1346
avg / total	0.86	0.86	0.85	3960

Confusion matrix:

```
[[ 521  112  224]
 [   3 1680   74]
 [   18  143 1185]]
```

Accuracy of GradientBoostingClassifier is 0.44772727272727275

Classification report for classifier GradientBoostingClassifier(criterion='friedman_mse',
init=None,

```
learning_rate=0.001, loss='deviance', max_depth=3,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=50,
presort='auto', random_state=None, subsample=1.0, verbose=0,
warm_start=False):
precision    recall  f1-score   support
```

negative	0.00	0.00	0.00	857
neutral	0.45	1.00	0.62	1757
positive	1.00	0.01	0.02	1346
avg / total	0.54	0.45	0.28	3960

Confusion matrix:

```
[[  0  857   0]
 [  0 1757   0]
 [  0 1330  16]]
```

In [23]:

```
print("Plotting the Model Performances: =====")
Index = [1,2,3,4,5,6,7,8,9]
plt.figure(1,figsize=(20, 10))
font = {'weight' : 'bold',
        'size' : 25}

plt.rc('font', **font)

plt.bar(Index,Accuracy)
plt.xticks(Index, Model,rotation=45)
plt.ylabel('Accuracy')
plt.xlabel('Model')
plt.title('Accuracies of Models')

print("Using the CountVectorizer on the Original Data: =====")
tweets = tweet_data['Text']
cv = CountVectorizer(ngram_range=(1,2), min_df=3, max_df=.95, stop_words='english')
bow = cv.fit_transform(tweets)

# use below if you need a data frame
bow_df = pd.DataFrame(bow.toarray(), index=tweets.index, columns=cv.get_feature_names())

X, Y = bow, (tweet_data['Polarity']).ravel()

binarize = Binarizer()
X = binarize.fit_transform(X)
```



```

X_train, X_test, y_train, y_test = \
    train_test_split(X, Y, test_size=0.3)

model = MultinomialNB()
model.fit(X_train, y_train)

preds = model.predict(X_test)

accuracy = accuracy_score(preds, y_test)
print("=====")
print("*****")
print('Accuracy is ' + str(accuracy))
print("-----")
print("Classification report for classifier %s:\n%s\n"
      % (classifier, metrics.classification_report(y_test, preds)))
cm = metrics.confusion_matrix(y_test, preds)
print("-----")
print("Confusion matrix:\n%s" % cm)
print("=====")

```

Plotting the Model Performances: =====

Using the CountVectorizer on the Original Data: =====

=====

Accuracy is 0.8058333333333333


Classification report for classifier GradientBoostingClassifier(criterion='friedman_mse',
init=None,

learning_rate=0.001, loss='deviance', max_depth=3,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=50,
presort='auto', random_state=None, subsample=1.0, verbose=0,
warm_start=False):

	precision	recall	f1-score	support
negative	0.78	0.66	0.72	749
neutral	0.96	0.78	0.86	1575
positive	0.70	0.92	0.79	1276
avg / total	0.83	0.81	0.81	3600

Confusion matrix:
[[498 15 236]
[71 1229 275]
[70 32 1174]]

=====



LogisticRegression

RandomForestClassifier

AdaBoostClassifier

BernoulliNB

MultinomialNB

KNeighborsClassifier

DecisionTreeClassifier

ExtraTreesClassifier

GradientBoostingClassifier

Model

In [24]:

```
X, Y = bow, (tweet_data['Polarity']).ravel()
ss = StandardScaler()
X = X.toarray()
X = ss.fit_transform(X)

X_train, X_test, y_train, y_test = \
    train_test_split(X, Y, test_size=0.3)

model = SVC()
model.fit(X_train, y_train)

preds = model.predict(X_test)

accuracy = accuracy_score(preds, y_test)
print("*****")
print('Accuracy is '+str(accuracy))
print("-----")
print("Classification report for classifier %s:\n%s\n"
      % (classifier, metrics.classification_report(y_test, preds)))
cm = metrics.confusion_matrix(y_test, preds)
print("-----")
print("Confusion matrix:\n%s" % cm)
print("=====")

models = [('mNB' , MultinomialNB()),
          ('bNB' , BernoulliNB()),
          ('svc' , SVC())]

print('{0}\t{1:<1}\t{2:<4}\t{3:<4}'.format("ACCURACY", "MEAN", "MIN", "MAX"))

*****
Accuracy is 0.8063888888888889
-----
Classification report for classifier GradientBoostingClassifier(criterion='friedman_mse',
init=None,
                      learning_rate=0.001, loss='deviance', max_depth=3,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=50,
                      presort='auto', random_state=None, subsample=1.0, verbose=0,
                      warm_start=False):
precision    recall  f1-score   support

negative     1.00      0.56      0.71       772
neutral      0.79      0.93      0.85      1580
positive     0.77      0.81      0.79      1248

avg / total     0.83      0.81      0.80      3600

-----
Confusion matrix:
[[ 429  149  194]
 [   0 1466  114]
 [   2  238 1008]]
=====
ACCURACY MEAN MIN  MAX
```

In [25]:

```

for name, model in models:
    X, Y = bow, (tweet_data['Polarity']).ravel()

    if name == 'bNB':
        binarize = Binarizer()
        X = binarize.fit_transform(X)
    elif name == 'svc':
        ss = StandardScaler()
        X = X.toarray()
        X = ss.fit_transform(X)

    cv = cross_val_score(model, X, Y, cv=5, scoring='accuracy')

    print('{0}\t{1:<3}\t{2:<4}\t{3:<4}'.format(name, round(cv.mean(), 4), round(cv.min(), 4), round(
cv.max(), 4)))

```

```

mNB 0.648 0.5427 0.7421
bNB 0.6791 0.5932 0.7621
svc 0.662 0.5685 0.7591

```

In [26]:

```

tweet_data.Text.head()

tweet_data.Text.apply(split_into_tokens)

bow_transformer = CountVectorizer(analyzer=split_into_lemmas).fit(tweet_data['Text'])
print("Printing the CountVectorizer BOW: =====")
print(len(bow_transformer.vocabulary_))
print("=====")

tweet1 = tweet_data['Text'][0]
print(tweet1)

bow1 = bow_transformer.transform([tweet1])
print(bow1)
print(bow1.shape)

print("=====")

```

```

Printing the CountVectorizer BOW: =====
40237
=====

```

```

89% of the cases reported today are asymptomatic.
If you are in a group of 10, assume nine of you have it!
Wear your mask properly and social distance.
Stay safe!

```

```

(0, 253) 1
(0, 1350) 1
(0, 1436) 1
(0, 2612) 1
(0, 3070) 2
(0, 3328) 1
(0, 3354) 1
(0, 5576) 1
(0, 9496) 1
(0, 13429) 1
(0, 13827) 1
(0, 14668) 1
(0, 14938) 1
(0, 15813) 1
(0, 18644) 1
(0, 20736) 1
(0, 21318) 3
(0, 23971) 1
(0, 25399) 1
(0, 26283) 1
(0, 27755) 1
(0, 28247) 1
(0, 34047) 1
(0, 34390) 1
(0, 36612) 1
(0, 37224) 2
(0, 37233) 1

```

```
(1, 40237)
```

```
In [27]:
```

```
tweets_bow = bow_transformer.transform(tweet_data['Text'])
print('sparse matrix shape:', tweets_bow.shape)
print('number of non-zeros:', tweets_bow.nnz)
print('sparsity: %.2f%%' % (100.0 * tweets_bow.nnz / (tweets_bow.shape[0] * tweets_bow.shape[1])))
print("=====")

tfidf_transformer = TfidfTransformer().fit(tweets_bow)
tfidf1 = tfidf_transformer.transform(bow1)
print("Printing the TF-IDF Vectors: =====")
print(tfidf1)
print("=====")
print("TF-IDF Shape =====")
print(tfidf1.shape)
print("=====")
tweets_tfidf = tfidf_transformer.transform(tweets_bow)
```

```
sparse matrix shape: (12000, 40237)
number of non-zeros: 321873
sparsity: 0.07%
```

```
=====
Printing the TF-IDF Vectors: =====
```

```
(0, 37233) 0.17023825244130142
(0, 37224) 0.269388204804755
(0, 36612) 0.19868285680668113
(0, 34390) 0.193090325525985
(0, 34047) 0.07384332794172017
(0, 28247) 0.20350308280208593
(0, 27755) 0.20734193860214528
(0, 26283) 0.21374339216488197
(0, 25399) 0.21042163368618824
(0, 23971) 0.23233463848108252
(0, 21318) 0.2821371783638204
(0, 20736) 0.23630146679872074
(0, 18644) 0.1607742858578646
(0, 15813) 0.11887685296562249
(0, 14938) 0.10133959242272426
(0, 14668) 0.15218409503709066
(0, 13827) 0.12138630362976884
(0, 13429) 0.2253684863238835
(0, 9496) 0.2272844955816938
(0, 5576) 0.15705174877315234
(0, 3354) 0.13231740875470657
(0, 3328) 0.23735418075323655
(0, 3070) 0.24490418752774623
(0, 2612) 0.09158969922734571
(0, 1436) 0.0821434755239757
(0, 1350) 0.21354060124203336
(0, 253) 0.20647724944119555
```

```
=====
TF-IDF Shape =====
```

```
(1, 40237)
```

```
In [28]:
```

```
print("*****")
print("Running Multinomial NB on TF-IDF")
print("*****")
polarity_detector = MultinomialNB().fit(tweets_tfidf, tweet_data['Polarity'])

print('predicted:', polarity_detector.predict(tfidf1[0]))
print('expected:', tweet_data.Polarity[0])

all_predictions = polarity_detector.predict(tweets_tfidf)
print(all_predictions)

print('accuracy', accuracy_score(tweet_data['Polarity'], all_predictions))
print('confusion matrix\n', confusion_matrix(tweet_data['Polarity'], all_predictions))
print('(row=expected, col=predicted)')
```

```
plt.figure(figsize=(10,10))
plt.matshow(confusion_matrix(tweet_data['Polarity'], all_predictions), cmap=plt.cm.binary,
interpolation='nearest')
plt.colorbar()
plt.ylabel('expected label')
plt.xlabel('predicted label')

print(classification_report(tweet_data['Polarity'], all_predictions))
print("=====")
```

```
*****
Running Multinomial NB on TF-IDF
*****
```

```
predicted: ['positive']
expected: positive
['positive' 'positive' 'positive' ... 'positive' 'positive' 'neutral']
accuracy 0.7835833333333333
confusion matrix
```


```
[[1233   70 1252]
 [  10 4122 1115]
 [  11  139 4048]]
```

```
(row=expected, col=predicted)
```

	precision	recall	f1-score	support
negative	0.98	0.48	0.65	2555
neutral	0.95	0.79	0.86	5247
positive	0.63	0.96	0.76	4198
avg / total	0.85	0.78	0.78	12000

```
=====
```

<Figure size 720x720 with 0 Axes>



In [29]:

```
print("*****")
print("Running Support Vector Machines on TF-IDF")
print("*****")
polarity_detector = SVC().fit(tweets_tfidf, tweet_data['Polarity'])
print('predicted:', polarity_detector.predict(tfidf1[0]))
print('expected:', tweet_data.Polarity[0])
```

```

all_predictions = polarity_detector.predict(tweets_tfidf)
print(all_predictions)

print('accuracy', accuracy_score(tweet_data['Polarity'], all_predictions))
print('confusion matrix\n', confusion_matrix(tweet_data['Polarity'], all_predictions))
print('(row=expected, col=predicted)')

plt.matshow(confusion_matrix(tweet_data['Polarity'], all_predictions), cmap=plt.cm.binary,
interpolation='nearest')
plt.colorbar()
plt.ylabel('expected label')
plt.xlabel('predicted label')

```


```

*****
Running Support Vector Machines on TF-IDF
*****
predicted: ['neutral']
expected: positive
['neutral' 'neutral' 'neutral' ... 'neutral' 'neutral' 'neutral']
accuracy 0.43725
confusion matrix
[[ 0 2555  0]
 [ 0 5247  0]
 [ 0 4198  0]]
(row=expected, col=predicted)

```

Out[29]:

Text(0.5,0,'predicted label')



In [30]:

```

print("=====")
print("Comparison Run: =====")

X_train, X_test, y_train, y_test = \
    train_test_split(tweet_data['Text'], tweet_data['Polarity'], test_size=0.2)

vectorizer = TfidfVectorizer(min_df=5, max_df = 0.8, sublinear_tf=True, use_idf=True, stop_words='english')
train_corpus_tfidf = vectorizer.fit_transform(X_train)

```

```
test_corpus_tf_idf = vectorizer.transform(x_test)

svm_model = LinearSVC()
nb_model = MultinomialNB()

svm_model.fit(train_corpus_tf_idf,y_train)
nb_model.fit(train_corpus_tf_idf,y_train)

svm_result = svm_model.predict(test_corpus_tf_idf)
nb_result = nb_model.predict(test_corpus_tf_idf)
```

=====
Comparison Run: =====


In [31]:

```
print('accuracy', accuracy_score(y_test, svm_result))
print('confusion matrix\n', confusion_matrix(y_test, svm_result))
print('(row=expected, col=predicted)')
plt.matshow(confusion_matrix(y_test, svm_result), cmap=plt.cm.binary, interpolation='nearest')
plt.colorbar()
plt.ylabel('expected label')
plt.xlabel('predicted label')
```

```
accuracy 0.87
confusion matrix
[[366  43  73]
 [ 20 986  55]
 [ 46  75 736]]
(row=expected, col=predicted)
```

Out[31]:

Text(0.5,0,'predicted label')



In [32]:


```
print('accuracy', accuracy_score(y_test, nb_result))
print('confusion matrix\n', confusion_matrix(y_test, nb_result))
print('(row=expected, col=predicted)')
plt.matshow(confusion_matrix(y_test, nb_result), cmap=plt.cm.binary, interpolation='nearest')
plt.colorbar()
```

```
plt.ylabel('expected label')
plt.xlabel('predicted label')
```

```
accuracy 0.7804166666666666
confusion matrix
[[246  14 222]
 [ 11 821 229]
 [ 18  33 806]]
(row=expected, col=predicted)
```

Out[32]:

```
Text(0.5,0,'predicted label')
```



In [33]:

```
clf = Pipeline([
    ('vectorizer', CountVectorizer(analyzer="word",
                                  tokenizer=word_tokenize,          # ! Comment line to include mark
                                  _negation and uncomment next line
                                  #tokenizer=lambda text: mark_negation(word_tokenize(text)),
                                  preprocessor=lambda text: text.replace("<br />", " "),
                                  max_features=10000) ),
    ('classifier', LinearSVC())
])

clf.fit(X_train, y_train)
clf.score(X_test, y_test)

clf = Pipeline([
    ('vectorizer', CountVectorizer(analyzer="word",
                                  #tokenizer=word_tokenize,          # ! Comment line to include mark
                                  k_negation and uncomment next line
                                  tokenizer=lambda text: mark_negation(word_tokenize(text)),
                                  preprocessor=lambda text: text.replace("<br />", " "),
                                  max_features=10000) ),
    ('classifier', LinearSVC())
])

clf.fit(X_train, y_train)
t = clf.score(X_test, y_test)
print("Analysis for the Linear SVC using CountVectorizer: ", t)
```


Analysis for the Linear SVC using CountVectorizer: 0.8508333333333333

In [34]:

```
print("=====")
print("Running the N-Grams on the Tweets: ")
print("=====")
bigram_clf = Pipeline([
    ('vectorizer', CountVectorizer(analyzer="word",
                                  ngram_range=(2, 2),
                                  tokenizer=word_tokenize,
                                  # tokenizer=lambda text: mark_negation(word_tokenize(text)),
                                  preprocessor=lambda text: text.replace("<br />", " "))),
    ('classifier', LinearSVC())
])

bigram_clf.fit(X_train, y_train)
t = bigram_clf.score(X_test, y_test)
print("Bi Gram Analysis Results with Ngram 2,2: ", t)
print("-----")
unigram_bigram_clf = Pipeline([
    ('vectorizer', CountVectorizer(analyzer="word",
                                  ngram_range=(1, 2),
                                  tokenizer=word_tokenize,
                                  # tokenizer=lambda text: mark_negation(word_tokenize(text)),
                                  preprocessor=lambda text: text.replace("<br />", " "))),
    ('classifier', LinearSVC())
])

unigram_bigram_clf.fit(X_train, y_train)
t = unigram_bigram_clf.score(X_test, y_test)
print("UNigram Bigram Analysis with ngram 1,2: ", t)
print("-----")
unigram_bigram_clf = Pipeline([
    ('vectorizer', CountVectorizer(analyzer="word",
                                  ngram_range=(1, 2),
                                  #tokenizer=word_tokenize,
                                  tokenizer=lambda text: mark_negation(word_tokenize(text)),
                                  preprocessor=lambda text: text.replace("<br />", " "))),
    ('classifier', LinearSVC())
])

unigram_bigram_clf.fit(X_train, y_train)
t = unigram_bigram_clf.score(X_test, y_test)
print("UNigram Bigram Analysis with ngram 1,2 and Tokennization: ", t)
print("-----")
print("END of the Analysis")
print("-----")
print("=====")
```

=====
Running the N-Grams on the Tweets:

=====
Bi Gram Analysis Results with Ngram 2,2: 0.80875

UNigram Bigram Analysis with ngram 1,2: 0.85

UNigram Bigram Analysis with ngram 1,2 and Tokennization: 0.84875

END of the Analysis

=====